

Pandas Lab Exercise (Kaggle Automobile Dataset)

We shall now test your skills in using Pandas package. We will be using the [automobiles Dataset](#) from Kaggle.

Answer each question asked below wrt the automobiles dataset. Load pandas as pd and upload the Automobile.csv file as auto

```
In [12]: import pandas as pd
```

Load the Automobile dataset into variable "auto"

```
In [13]: auto=pd.read_csv("Automobile.csv")
         type(auto)
```

```
Out[13]: pandas.core.frame.DataFrame
```

Check the head of the DataFrame.

```
In [14]: auto.head()
```

```
Out[14]:
```

	symboling	normalized_losses	make	fuel_type	aspiration	number_of_doors	body_s
0	3	168	alfa-romero	gas	std	two	conver
1	3	168	alfa-romero	gas	std	two	conver
2	1	168	alfa-romero	gas	std	two	hatchl
3	2	164	audi	gas	std	four	se
4	2	164	audi	gas	std	four	se

5 rows × 26 columns

How many rows and columns are there?

```
In [15]: auto.shape
```

```
Out[15]: (201, 26)
```

```
In [ ]:
```

What is the average Price of all cars in the dataset?

```
In [16]: auto['price'].mean()
```

Out[16]: np.float64(13207.129353233831)

Which is the cheapest make and costliest make of car in the lot?

In [20]: `auto['price'].max() , auto['price'].idxmax()`

Out[20]: (np.int64(45400), 71)

In [21]: `auto['price'].min(), auto['price'].idxmin()`

Out[21]: (np.int64(5118), 134)

How many cars have horsepower greater than 100?

In [22]: `horse_power=auto[auto['horsepower']>100]
horse_power`

Out[22]:

	symboling	normalized_losses	make	fuel_type	aspiration	number_of_doors	body
0	3	168	alfa-romero	gas	std	two	conv
1	3	168	alfa-romero	gas	std	two	conv
2	1	168	alfa-romero	gas	std	two	hat
3	2	164	audi	gas	std	four	
4	2	164	audi	gas	std	four	
...	
196	-1	95	volvo	gas	std	four	
197	-1	95	volvo	gas	turbo	four	
198	-1	95	volvo	gas	std	four	
199	-1	95	volvo	diesel	turbo	four	
200	-1	95	volvo	gas	turbo	four	

90 rows × 26 columns

How many hatchback cars are in the dataset ?

In [23]: `hatchback_count=auto[auto['body_style']=='hatchback'].shape[0]`

In [24]: `hatchback_count`

Out[24]: 68

What are the 3 most commonly found cars in the dataset?

```
In [25]: most_common_cars=auto['make'].value_counts().head(3)
```

```
In [26]: most_common_cars
```

```
Out[26]: make
toyota    32
nissan     18
mazda     17
Name: count, dtype: int64
```

Someone purchased a car for 7099, what is the make of the car?

```
In [27]: make_Car=auto[auto['price']==7099]['make']
make_Car
```

```
Out[27]: 87    nissan
Name: make, dtype: object
```

Which cars are priced greater than 40000?

```
In [28]: greater_than_40k=auto[auto['price']>40000][['make','price']]
greater_than_40k
```

```
Out[28]:
```

	make	price
15	bmw	41315
70	mercedes-benz	40960
71	mercedes-benz	45400

Which are the cars that are both a sedan and priced less than 7000?

```
In [30]: cond = auto[(auto['body_style'] == 'sedan') & (auto['price'] < 7000)][['body_styl
cond
```

```
Out[30]:
```

	body_style
19	sedan
24	sedan
42	sedan
50	sedan
82	sedan
86	sedan
88	sedan
89	sedan
118	sedan
152	sedan

Count the number of unique values in the fuel_type column.

```
In [31]: unique_vals=auto['fuel_type'].unique()  
unique_vals
```

```
Out[31]: array(['gas', 'diesel'], dtype=object)
```

List all the cars that have a horsepower between 100 and 200, and display their make , horsepower , and price .

```
In [33]: cond_2= auto[(auto['horsepower']>100) & (auto['horsepower']<=200) ]  
cond_2[['make','horsepower','price']]
```

Out[33]:

	make	horsepower	price
0	alfa-romero	111	13495
1	alfa-romero	111	16500
2	alfa-romero	154	16500
3	audi	102	13950
4	audi	115	17450
...
196	volvo	114	16845
197	volvo	160	19045
198	volvo	134	21485
199	volvo	106	22470
200	volvo	114	22625

86 rows × 3 columns

Find the average city_mpg and highway_mpg for each body_style .

```
In [34]: avg_mpg = auto.groupby('body_style')[['city_mpg','highway_mpg']].mean()  
avg_mpg
```

Out[34]:

	city_mpg	highway_mpg
body_style		
convertible	20.500000	26.000000
hardtop	21.625000	27.250000
hatchback	26.602941	32.382353
sedan	25.053191	30.574468
wagon	24.040000	28.720000

What is the median price for each make ?

```
In [35]: median_price = auto.groupby('make')['price'].mean()  
median_price
```

```
Out[35]: make
         alfa-romero      15498.333333
         audi           17859.166667
         bmw            26118.750000
         chevrolet       6007.000000
         dodge          7875.444444
         honda          8184.692308
         isuzu           8916.500000
         jaguar          34600.000000
         mazda          10652.882353
         mercedes-benz   33647.000000
         mercury        16503.000000
         mitsubishi      9239.769231
         nissan          10415.666667
         peugot         15489.090909
         plymouth        7963.428571
         porsche         31400.500000
         renault         9595.000000
         saab           15223.333333
         subaru          8541.250000
         toyota          9885.812500
         volkswagen     10077.500000
         volvo          18063.181818
         Name: price, dtype: float64
```

List all cars that have a `wheel_base` greater than 100 and a `curb_weight` less than 2500.

```
In [36]: cond_3 = auto[(auto['wheel_base'] > 100) & (auto['curb_weight'] < 2500)]
         cond_3
```

```
Out[36]:
```

	symboling	normalized_losses	make	fuel_type	aspiration	number_of_doors	body
9	2	192	bmw	gas	std	two	
10	0	192	bmw	gas	std	four	
169	-1	65	toyota	gas	std	four	
170	-1	65	toyota	diesel	turbo	four	
171	-1	65	toyota	gas	std	four	hatc
172	-1	65	toyota	gas	std	four	
173	-1	65	toyota	gas	std	four	hatc

7 rows × 26 columns

Create a new column `price_per_hp` that calculates the price of the car per horsepower.

```
In [37]: auto['price_per_hp'] = auto['price'] / auto['horsepower']
         auto.head()
```

```
Out[37]:
```

	symboling	normalized_losses	make	fuel_type	aspiration	number_of_doors	body_s
0	3	168	alfa-romero	gas	std	two	conver
1	3	168	alfa-romero	gas	std	two	conver
2	1	168	alfa-romero	gas	std	two	hatchl
3	2	164	audi	gas	std	four	se
4	2	164	audi	gas	std	four	se

5 rows × 27 columns

Count how many cars have a `number_of_doors` as `four` .

```
In [41]: four_nofdoors = auto[auto['number_of_doors'] == 'four'].shape[0]
four_nofdoors
```

```
Out[41]: 114
```

Find the top 5 cars based on their `highway_mpg` and `price` .

```
In [42]: top_5_cars=auto.sort_values(['highway_mpg','price'],ascending=[False,False]).head(5)
top_5_cars
```

```
Out[42]:
```

	symboling	normalized_losses	make	fuel_type	aspiration	number_of_doors	bo
29	2	137	honda	gas	std	two	h
17	2	121	chevrolet	gas	std	two	h
87	1	128	nissan	diesel	std	two	
155	0	91	toyota	diesel	std	four	h
156	0	91	toyota	gas	std	four	

5 rows × 27 columns

How many cars have missing values in the `normalized_losses` column?

```
In [43]: missing_vals = auto['normalized_losses'].isnull().sum()
missing_vals
```

```
Out[43]: np.int64(0)
```

Create a new column `car_age` that calculates the age of the car based on the `year_of_manufacture` (assume the current year is 2025).

```
In [46]: current_year = 2025
auto['car_age'] = current_year - auto['year_of_manufacture']
auto.head()
```

```
-----
KeyError                                Traceback (most recent call last)
File C:\Python311\Lib\site-packages\pandas\core\indexes\base.py:3805, in Index.get_loc(self, key)
    3804 try:
-> 3805     return self._engine.get_loc(casted_key)
    3806 except KeyError as err:

File index.pyx:167, in pandas._libs.index.IndexEngine.get_loc()

File index.pyx:196, in pandas._libs.index.IndexEngine.get_loc()

File pandas\_libs\hashtable_class_helper.pxi:7081, in pandas._libs.hashtable.PyObjectHashTable.get_item()

File pandas\_libs\hashtable_class_helper.pxi:7089, in pandas._libs.hashtable.PyObjectHashTable.get_item()
```

KeyError: 'year_of_manufacture'

The above exception was the direct cause of the following exception:

```
KeyError                                Traceback (most recent call last)
Cell In[46], line 2
      1 current_year = 2025
----> 2 auto['car_age'] = current_year - auto['year_of_manufacture']
      3 auto.head()

File C:\Python311\Lib\site-packages\pandas\core\frame.py:4102, in DataFrame.__getitem__(self, key)
    4100 if self.columns.nlevels > 1:
    4101     return self._getitem_multilevel(key)
-> 4102 indexer = self.columns.get_loc(key)
    4103 if is_integer(indexer):
    4104     indexer = [indexer]

File C:\Python311\Lib\site-packages\pandas\core\indexes\base.py:3812, in Index.get_loc(self, key)
    3807 if isinstance(casted_key, slice) or (
    3808     isinstance(casted_key, abc.Iterable)
    3809     and any(isinstance(x, slice) for x in casted_key)
    3810 ):
    3811     raise InvalidIndexError(key)
-> 3812     raise KeyError(key) from err
    3813 except TypeError:
    3814     # If we have a listlike key, _check_indexing_error will raise
    3815     # InvalidIndexError. Otherwise we fall through and re-raise
    3816     # the TypeError.
    3817     self._check_indexing_error(key)
```

KeyError: 'year_of_manufacture'

The END